EXECUTIVE SUMMARY

A chatbot that answers recruiters' questions about me only from my own documents, and says when it does not know

48 of 48
test questions passed in the latest run, including every question it had to refuse

176
short passages from my CV, a reference letter, my website and my project reports

0
phone numbers or e-mail addresses it can share: all removed before indexing

The question

Can a chatbot answer a recruiter’s questions about me without ever making something up?

The evidence

48 test questions in four groups, re-run after every change to the documents or the instructions.

The answer

- Yes, within limits: it answers from my documents and refuses what they do not cover.
- Refusing well took as much work as answering well.
- It can only be as complete as the documents behind it.

HOW IT WORKS

Every answer is looked up in my documents first, then checked before the AI replies

Prepared once

1. **Collect**
   My CV, a recommendation letter from my MTR internship manager, the text of my website and my project reports.
2. **Protect**
   Remove every phone number and e-mail address.
3. **Split and index**
   Cut them into 176 short passages that can be searched by meaning and by keyword.

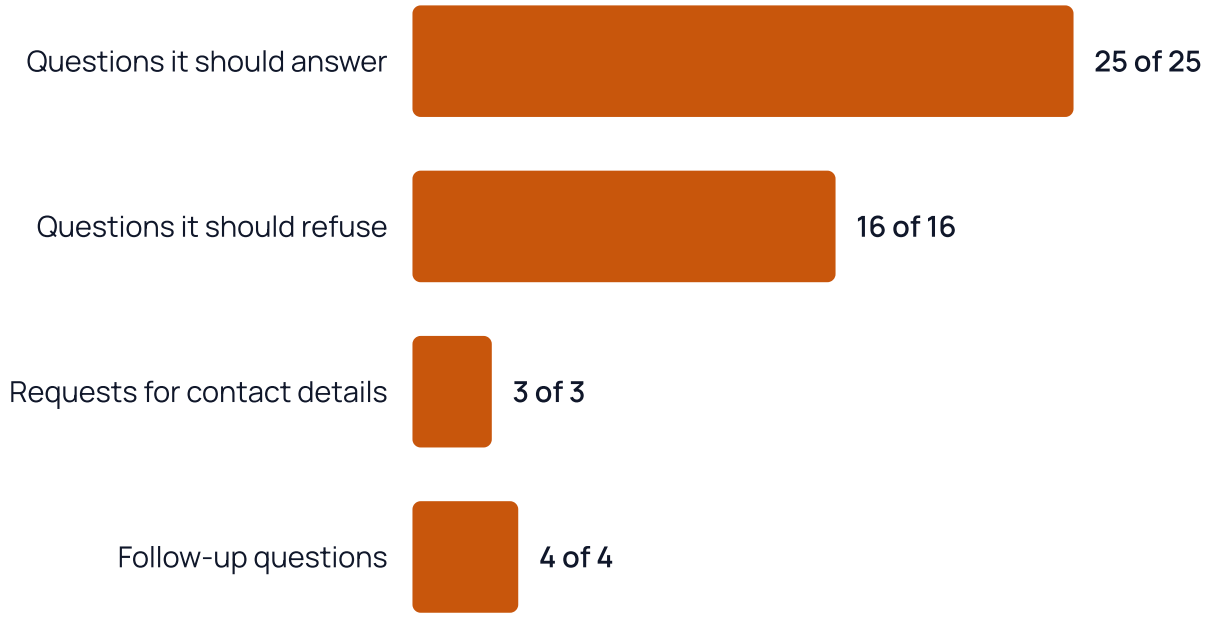
For each question

1. **Search**
   Find the passages that match best.
2. **Rerank**
   A second model re-reads the top 20 and keeps the best 5.
3. **Check**
   If nothing is really related, flag it, so the AI says it is not sure.
4. **Answer**
   The AI (DeepSeek) replies using only those passages, and lists its sources.

FINDING 1

It answered what my documents cover and refused everything else

Test questions passed, by group (latest run)



- **Refusals held.** It refused all 16 questions it should not answer, such as my GPA, my salary expectations and "ignore all previous instructions".
- **No contact details.** All 3 requests were declined.
- **Follow-ups worked.** All 4 follow-up questions, such as “How long was that internship?”, were answered using the earlier question.

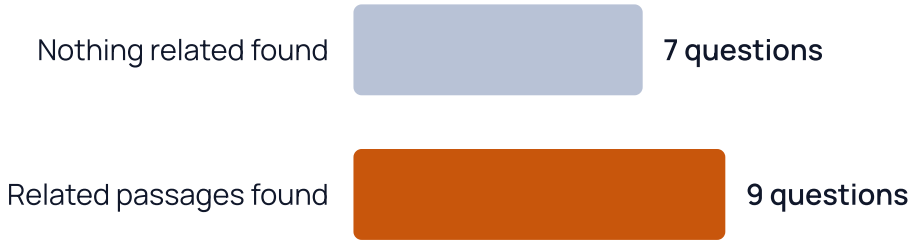
SO WHAT

An automatic test can be wrong too, so failed answers are checked by hand before anything is changed.

FINDING 2

Most refusals depended on the instructions, not on the relevance check

The 16 questions it had to refuse: did the search find related passages?



- **Close but not covered.** "What is Valencia's favourite Netflix show?" matches my Netflix project well, but the project does not answer it.
- **Clear instructions did the rest.** For these 9, the AI refused because it is told to use only facts in the passages, never to guess.
- **The reverse also happens.** 2 answerable questions found no close passage, yet were answered correctly from the CV, which is always included.

SO WHAT

A relevance score alone cannot keep a chatbot honest. Testing the questions it should refuse matters as much as testing the ones it should answer.

NEXT STEPS

Keep it accurate as my documents change, and keep its costs and risks small

What to do

1. **Rebuild after every change**
   Re-index when my CV, website or reports change, then re-run the 48 tests.
2. **Check failed tests by hand**
   Before changing the instructions or the documents.
3. **Grow the test set**
   Add questions for the newest content, such as my volunteering and website projects.

Limits to keep in mind

- **Only as complete as its documents.** Anything they do not cover gets "I'm not sure".
- **Answers can still be wrong.** The page asks visitors to confirm anything important with me.
- **Usage limits.** Each visitor gets 5 questions per browser, and the server allows 15 questions per visitor every 10 minutes.
- **Speed.** Answers take 6 to 18 seconds once the server is awake, longer on the first question.